



Apache Flume

Getting data into Hadoop

Problem

- Getting data into HDFS is not difficult:
 - **% `hadoop fs --put data.csv` .**
 - works great when data is neatly packaged and ready to upload
- Unfortunately, e.g. a webserver is creating data all the time
 - How often should a batch load data to HDFS happen? Daily? Hourly?
- The real need is a solution that can deal with streaming logs/data



Solution: Apache Flume

- Introduced in Cloudera's **CDH3** distribution
- versions 0.x: flume, 1.x: flume-ng



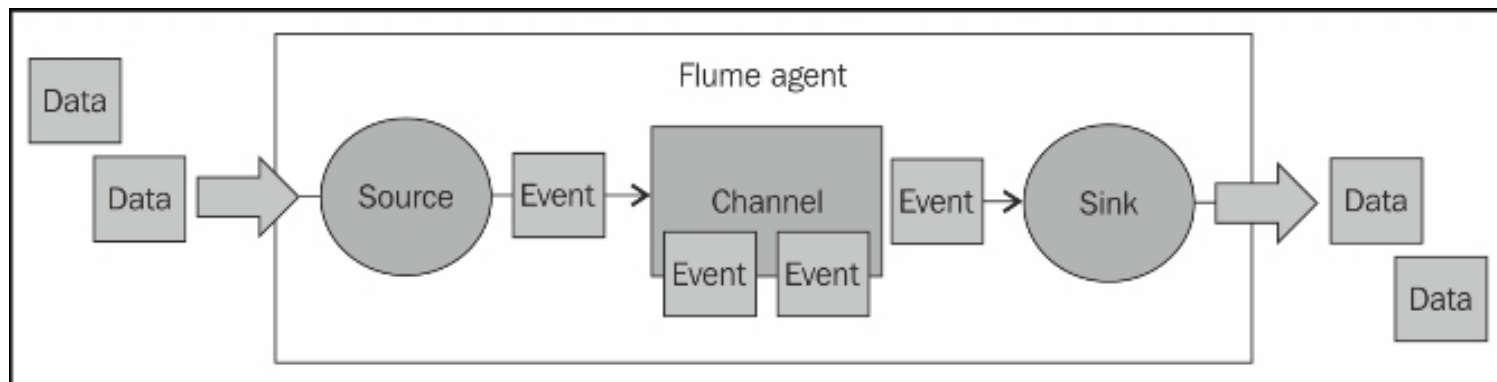
Overview

- Stream data (events, not files) from clients to sinks
- Clients: files, syslog, avro, ...
- Sinks: HDFS files, HBase, ...
- Configurable reliability levels
 - Best effort: “Fast and loose”
 - Guaranteed delivery: “Deliver no matter what”
- Configurable routing / topology



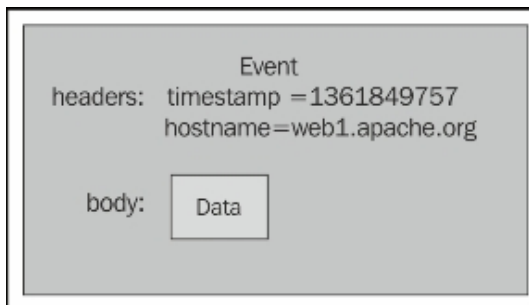
Architecture

Component	Function
Agent	The JVM running Flume. One per machine. Runs many sources and sinks.
Client	Produces data in the form of events. Runs in a separate thread.
Sink	Receives events from a channel. Runs in a separate thread.
Channel	Connects sources to sinks (like a queue). Implements the reliability semantics.
Event	A single datum; a log record, an avro object, etc. Normally around ~4KB.



Events

- Payload of the data is called an *event*
 - composed of zero or more headers and a body
- Headers are key/value pairs
 - making routing decisions or
 - carry other structured information



Channels

- Provides a buffer for in-flight events
 - after they are read from sources
 - until they can be written to sinks in the data processing pipelines
- Two (three) primary types are
 - a memory-backed/nondurable channel
 - a local-filesystem-backed/durable channel
 - (hybrid)



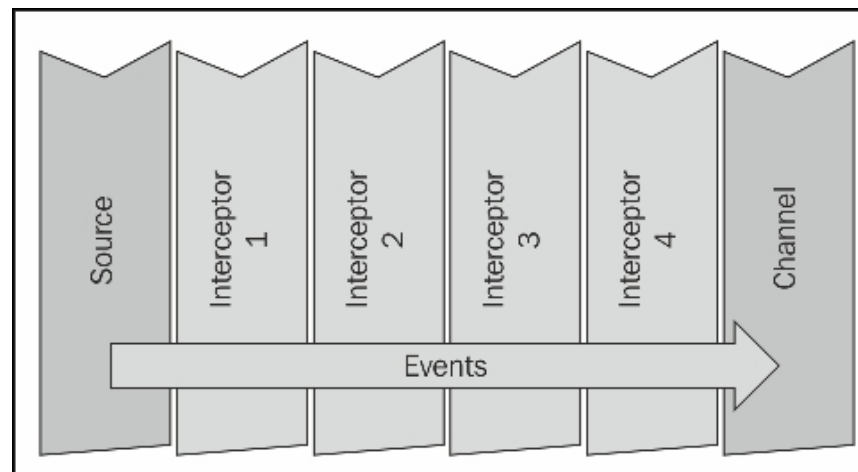
Channels

- The writing rate of the sink should be faster than the ingest rate from the sources
 - `ChannelException` might lead to data loss



Interceptors

- An **interceptor** is a point in data flow where events can be inspected and altered
- zero or more interceptors can be chained after a source creates an event



Channel Selectors

- Responsible for how data moves from a source to one or more channels
- Flume comes with two selectors
 - **replicating channel selector** (the default) puts a copy of the event into each channel
 - **multiplexing channel selector** writes to different channels depending on headers
- Combined with interceptors forms the foundation for routing



Sinks

- Flume supports a set of sinks
 - HDFS, ElasticSearch, Solr, HBase, IRC, MongoDB, Cassandra, RabbitMQ, Redis, ...
- HDFS Sink continuously
 - open a file in HDFS,
 - stream data into it,
 - at some point, close that file
 - start a new one

`agent.sinks.k1.type=hdfs`

`agent.sinks.k1.hdfs.path=/path/in/hdfs`



Sources

- Flume source consumes events delivered to it by an external source
 - like a web server



Tiered Collection

- Send events from agents to another tier of agents to aggregate
- Use an Avro sink (really just a client) to send events to an Avro source (really just a server) in another machine
- Failover supported
- Load balancing (soon)
- Transactions guarantee handoff

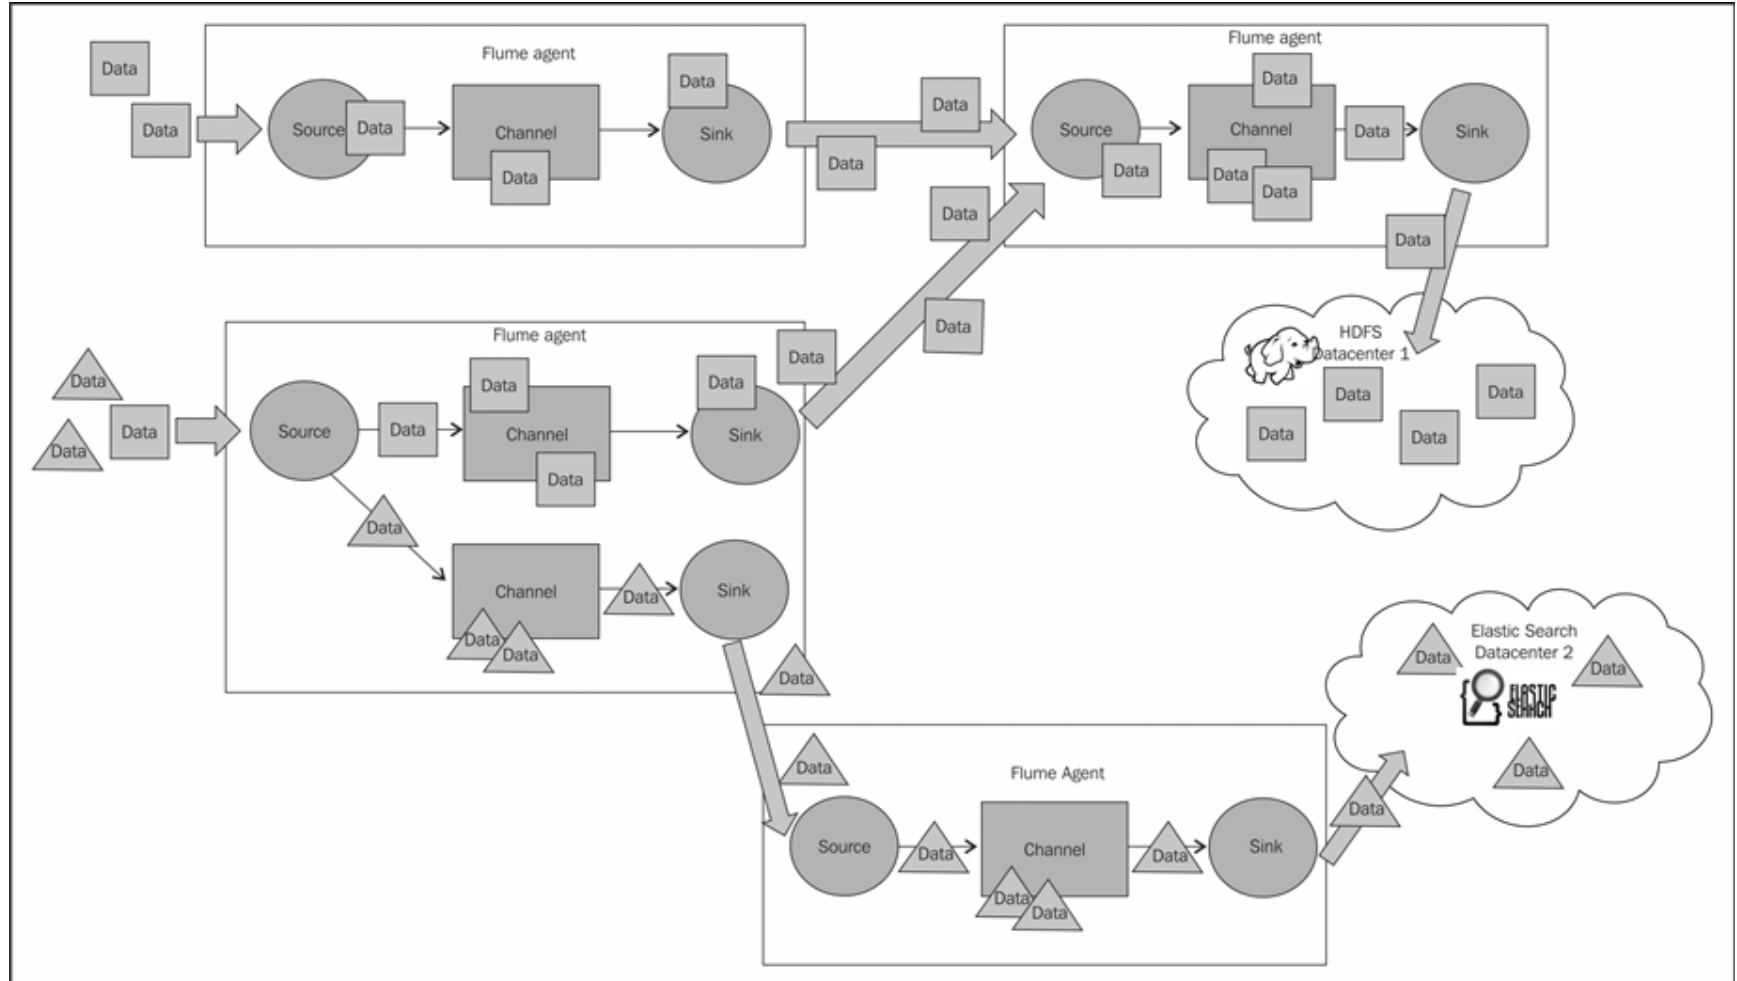


Tiered Collection Handoff

- Agent 1: Tx begin
- Agent 1: Channel take event
- Agent 1: Sink send
- Agent 2: Tx begin
- Agent 2: Channel put
- Agent 2: Tx commit, respond OK
- Agent 1: Tx commit (or rollback)



Tiered Data Collection



Apache Flume

- A source writes events to one or more channels
- A channel is the holding area as events are passed from a source to a sink
- A sink receives events from one channel only
- An agent can have many channels



Flume Configuration File

- Simple Java property file of key/value pairs
- Several agents can be configured in a single file
 - agents are identified by **agent identifier** (called a **name**)
- Each agent is configured, starting with three parameters:

`agent.sources=<list of sources>`

`agent.channels=<list of channels>`

`agent.sinks=<list of sinks>`



- Each source, channel and sink has a unique name within the context of that agent
 - Prefix for channel named access
`agent.channels.access`
- Each item has a type
 - E.g. in-memory channel is *memory*
`agent.channels.access.type=memory`



Hello, World!

```
agent.sources=s1  
agent.channels=c1  
agent.sinks=k1
```

```
agent.sources.s1.type = spooldir  
agent.sources.s1.spoolDir = /etc/spool  
...  
agent.sinks.k1.type = hdfs  
agent.sinks.k1.hdfs.path =  
hdfs://localhost:9001/user/hduser/log-data  
...
```



Hello, World!

- Config has one agent (called *agent*) with
 - a source named *s1*
 - a channel named *c1*
 - a sink named *k1*
- The *s1* source's type is *spooldir*
 - Files appearing in */etc/spool* are ingested
- The type of the sink named *k1* is *hdfs*
 - writes data files to *log-data*



Command Line Usage

% flume-ng help

Usage: /usr/local/flume/apache-flume-1.6.0-bin/bin/flume-ng <command> [options]...

commands:

help

display this help text

agent

run a Flume agent

avro-client

run an avro Flume client

version

show Flume version info

...



Command Line Usage

- The agent command requires 2 parameters
 - a configuration file to use and
 - the agent name
- Example

```
% flume-ng agent -n agent -f myConf.conf ...
```
- Test

```
% cp log-data/* /etc/spool
```



Sources Code

```
public class MySource implements PollableSource {
    public Status process() {
        // Do something to create an Event..
        Event e = EventBuilder.withBody(...).build();
        // A channel instance is injected by Flume.
        Transaction tx = channel.getTransaction();
        tx.begin();
        try {
            channel.put(e);
            tx.commit();
        } catch (ChannelException ex) {
            tx.rollback();
            return Status.BACKOFF;
        } finally {
            tx.close();
        }
        return Status.READY;
    }
}
```



Sinks Code

```
public class MySink implements PollableSink {
    public Status process() {
        Transaction tx = channel.getTransaction();
        tx.begin();
        try {
            Event e = channel.take();
            if (e != null) {
                // ...
                tx.commit();
            } else {
                return Status.BACKOFF;
            }
        } catch (ChannelException ex) {
            tx.rollback();
            return Status.BACKOFF;
        } finally {
            tx.close();
        }

        return Status.READY;
    }
}
```

